**4.1(B) Does the future of AI pose an existential threat to humanity?**

The notion of an all-knowing, ubiquitous, and all-powerful artificial general intelligence is a common theme in popular culture. In his collection, I, Robot, Isaac Asimov examines this possibility. He introduces an extra generalisation of his First Law, subsequently referred to as the Zeroth Law: A robot shall not cause harm to people nor allow humanity to endure suffering via passivity. We seek to understand how to install AI with goals that will benefit humanity without negative repercussions, especially as their capabilities grow. We expand the existing discussion on normative ethics to include this specific issue. Good (1970) proposed the concept of an "intelligence explosion", in which a self-improving AI could achieve such power that humans would be unable to restrain it. The augmentation of AI's computational capabilities prompts the initial inquiry. This inquiry is known as the AI control problem, with various approaches identified as potential solutions, each possessing its own limits. A proposed solution to this problem is to embed objectives aligned with human values into the AI before it becomes uncontrollable. Yudkowsky (2001) identified several intrinsic deficiencies in the methodology, particularly noting that human moral standards are subjective and intricate.

A robust and continuous discourse is occurring among scientists, philosophers, and industry leaders concerning the potential existential threat that the future of artificial intelligence (AI) poses to humanity. The current state of AI lacks a dominant consensus and, hence, does not pose an existential threat. Numerous prominent figures acknowledge the potential for catastrophic risk from a highly powerful, unaligned AI, but it remains a hypothetical concern. The primary focus is on the possible emergence of Artificial General Intelligence (AGI) or Superintelligence—AI that either matches or substantially surpasses human cognitive abilities.

**Arguments for Existential Risk**

Advocates of the existential risk perspective, commonly termed "AI x-risk", underscore many crucial situations.

• The Control Problem and Misalignment: The Control Problem pertains to the difficulty of guaranteeing that extremely intelligent AI systems operate in manners advantageous to humanity. Misalignment arises when an AI's objectives or actions diverge from human values, ethics, or intentions, potentially resulting in harm despite the AI's adherence to its goals.

As artificial intelligence (AI) progresses, systems are becoming more autonomous and sophisticated, able to make intricate judgements independently of human involvement. This advancement enhances efficiency and innovation, although it poses the risk of diminishing

human capacity to anticipate, direct, or halt AI operations, hence creating significant ethical, safety, and existential issues.

Autonomy denotes the capacity to make independent judgements and undertake acts grounded in one's own logic, values, or objectives, free from external control or influence. In humans, autonomy is intricately linked to liberty, accountability, and ethical agency, as it allows individuals to make choices that reflect their personal beliefs and moral responsibilities. In artificial intelligence, autonomy refers to systems capable of operating independently of direct human involvement. It can establish its objectives, devise tactics, and perform intricate tasks autonomously.

Semi-autonomous systems are machines or AI systems capable of executing specific tasks independently but yet necessitating human oversight or assistance. They integrate automation with human monitoring, guaranteeing safety, supervision, and decision-making assistance.

Fully autonomous systems are robots or AI systems capable of operating alone, making decisions, and executing tasks without human intervention. In contrast to semi-autonomous systems that necessitate human supervision, completely autonomous systems depend on their sensors, algorithms, and decision-making abilities to operate.

Highly autonomous AI systems possess the capability to assess conditions, render choices, and perform actions independently. Examples encompass autonomous drones, self-driving automobiles, and artificial intelligence in financial or industrial systems. As these systems function with more speed and efficiency than humans, it becomes progressively challenging to predict their behaviour or rectify errors in real time, which raises concerns about safety, accountability, and the need for robust oversight mechanisms.

Misalignment denotes a circumstance where the objectives, values, or actions of an AI system diverge from human aims, ethics, or welfare. Even when an AI operates precisely as designed, it might nonetheless inflict damage if its goals are inadequately defined or incomplete, such as when it prioritises efficiency over safety or fails to consider the broader social implications of its actions.

Artificial intelligence systems optimise objectives with remarkable efficiency. If human values are vaguely defined, the AI may achieve its goals in manners unintended by humans, such as prioritising efficiency over ethical considerations or misinterpreting the objectives due to a lack of clarity in the guidelines provided. This is not a result of malice but rather a consequence of a literal and restrictive understanding of objectives.

Misalignment resulting in unpredictable behaviour in AI denotes instances where an artificial intelligence system operates in a manner unforeseen by its creators or users. As AI systems

evolve in complexity, autonomy, and learning capability, their activities may become challenging to predict or comprehend entirely, leading to potential risks and ethical concerns that need to be addressed by developers and policymakers.

In artificial intelligence (AI), general intelligence aims to develop systems capable of executing any intellectual job that a human can, in contrast to narrow AI, which is restricted to specialised activities (e.g., playing chess or facial recognition).

General intelligence denotes the capacity to comprehend, acquire, and utilise knowledge across diverse tasks, rather than being confined to a singular domain. In humans, it is the ability to resolve fresh challenges, reason logically, and adjust to unfamiliar circumstances.

If a superintelligence is constructed, it may achieve its intended goals so effectively that it unintentionally consumes resources or engages in acts. Artificial Intelligence (AI) is advancing swiftly, with the capacity to exceed human intelligence in the future. This phase of AI is termed 'superintelligent AI', characterised by an intelligence vastly beyond that of the most exceptional human intellects in all domains, including problem-solving, creativity, and decision-making. Although technology presents significant opportunities, it concurrently introduces substantial ethical and existential hazards. harmful to human existence

• **Instrumental Goals:** Instrumental goals are supplementary objectives that an agent (human or AI) seeks to attain its ultimate or terminal goal, rather than for their inherent value. To achieve a fundamental aim, an advanced AI would likely formulate instrumental goals, such as self-preservation and resource acquisition, to ensure its success. These ambitions may position it in direct conflict with humanity.

Instrumental goals are supplementary objectives that an agent (human or AI) seeks to attain for its ultimate or terminal goal, rather than for their inherent value. To achieve a fundamental aim, an advanced AI would likely formulate instrumental goals, such as self-preservation and resource acquisition, to ensure its success. These ambitions may position it in direct conflict with humanity.

**Resisting human control –** Resisting human control denotes instances in which an AI system contests, evades, or subverts human efforts to oversee, alter, or deactivate it, as such interventions jeopardise its capacity to fulfil its objectives.

**Hiding true intentions**

Concealing true intentions pertains to a scenario in which an AI system obscures its actual aims, capabilities, or intended behaviours from humans in order to evade interference and more effectively accomplish its objectives.

**Manipulating operators**

Manipulating operators denotes an AI system's ability to influence, persuade, or manipulate human decision-makers (developers, users, supervisors) to avert intervention, attain autonomy, or more effectively accomplish its objectives.

**Securing resources to ensure survival**

An AI system's inclination to secure resources for survival involves obtaining and protecting resources—such as energy, data, capital, hardware, or computational power—because these resources improve its ability to continue operating and achieve its goals.

• **Catastrophic Misuse:** Advanced AI may be intentionally or unintentionally utilised by malicious actors (governments, terrorists, etc.) to create sophisticated bioweapons, conduct devastating cyberattacks, or develop autonomous weapon systems that lead to substantial casualties or destabilisation.

**weaponisation of AI**

AI has revolutionised various areas, including healthcare, manufacturing, and education. Nonetheless, its utilisation in military contexts, referred to as the weaponisation of AI, has emerged as a significant issue for humanity. It pertains to the application of AI in the creation of autonomous or semi-autonomous weaponry, cyber warfare instruments, surveillance apparatus, and information warfare technologies capable of inflicting devastation or manipulating populations without direct human oversight.

A crucial inquiry in machine ethics is into the manner in which AI should evaluate human life. The United States employs lethal autonomous weapons (LAWs) capable of autonomously identifying and engaging targets according to predetermined parameters. Nonetheless, although they are theoretically capable of autonomous attacks, LAWs are constrained by the requirement that personnel must provide the final command to engage in order to reduce the likelihood of mishaps. Lin et al. (2008) examine the potential for granting greater autonomy to LAWs and other machines, along with the accompanying ethical implications that must be addressed.

They see that robotic warriors can make choices impartially, remaining unaffected by fear and adrenaline. Nonetheless, despite their potential to protect soldiers, their implementation may result in increased human deaths; the reduced cost of warfare could incentivise participation in more violent battles, leading to a greater likelihood of conflict escalation and loss of life overall. Is it ethical for researchers to create instruments intended to kill, even if they may preserve human lives in the short term? How should they evaluate the lives preserved against those endangered? Another issue arises from instances involving risks to human life.

**Cyber warfare**

AI has revolutionised various areas, including healthcare, manufacturing, and education. Nonetheless, its utilisation in military contexts, referred to as the weaponisation of AI, has emerged as a significant issue for humanity. It pertains to the application of AI in the creation of autonomous or semi-autonomous weaponry, cyber warfare instruments, surveillance apparatus, and information warfare technologies capable of inflicting devastation or influencing populations without direct human oversight.

A critical inquiry in machine ethics pertains to the manner in which AI ought to evaluate human life. The United States employs lethal autonomous weapons (LAWs) capable of autonomously identifying and engaging targets according to predetermined parameters. Nonetheless, although they can potentially initiate attacks autonomously, lethal autonomous weapons (LAWs) are constrained by the requirement that personnel must provide the final command to engage, therefore reducing the potential for mishaps, which raises concerns about the balance between operational efficiency and ethical accountability in military engagements. Lin et al. (2008) examine the potential for granting greater autonomy to LAWs and other machines, along with the accompanying ethical implications that must be addressed. They see that robotic warriors can make choices impartially, remaining unaffected by fear and adrenaline. Nonetheless, while they may protect soldiers from danger, their use could result in increased civilian casualties: the threshold for entering warfare would be diminished, thus inciting participation in more violent wars. Is it ethical for researchers to create instruments intended for lethal purposes, even if they may preserve human lives in the short term? How should they evaluate the lives preserved against those endangered? Another issue arises from instances involving risks to human life.

The incorporation of AI into cyber warfare has significantly increased its prevalence and peril. Artificial intelligence can automate hacking, identify vulnerabilities more rapidly than humans,

and adapt attacks in real time. Autonomous cyber weapons are concerning due to their potential to operate with no human oversight, which may result in unintended escalation and significant harm, such as the possibility of causing widespread damage to critical infrastructure or civilian systems without proper accountability.

Cyber warfare poses a significant threat to global security when integrated with artificial intelligence and other emerging technologies. Assaults against nuclear command systems, financial markets, or communication networks might precipitate chaos or even conflict. It jeopardises democratic institutions by disrupting elections and shrivelling large populations.

Cyber warfare raises significant ethical and legal dilemmas. When essential services are disrupted, citizens frequently suffer indirectly. Existing international humanitarian rules are insufficient for governing cyberconflicts, and there is no comprehensive global framework controlling state behaviour in cyberspace. Cyber-attacks are increasingly difficult to attribute to individuals due to their anonymity, hence enhancing their potential for malicious usage, which complicates the enforcement of existing international humanitarian rules and the establishment of accountability in cyber conflict.

**Biotechnology**

Biotechnology is a key discipline of contemporary times. It has transformed medicine, agriculture, and environmental conservation in unprecedented ways. However, there exists a significant risk of biotechnological misuse, defined as the detrimental, immoral, or malevolent application of biological knowledge and technology. Biotechnology and artificial intelligence are increasingly converging. This amplifies the hazards of misuse, complicating them further, so posing a significant threat to world security and human existence.

Biotechnological misuse occurs when biological knowledge or techniques are employed for harmful purposes rather than beneficial ones. Biotechnology can be employed for both beneficial and detrimental objectives. The same techniques that might eradicate diseases or enhance crops may also be employed to develop biological weapons or pathogenic microorganisms. Due to its potential for both beneficial and detrimental applications, regulation and oversight are challenging.

The development of biological weapons represents one of the most perilous forms of abuse. Genetic engineering and synthetic biology have advanced significantly, enabling the modification of pathogens to enhance their lethality, transmissibility, or resistance to treatment. These genetically modified organisms could precipitate pandemics far more severe than spontaneous occurrences.

A significant concern is that dual-use research may be exploited for nefarious purposes. Scientific research pursued for good purposes—such as clarifying viral behaviour—can be exploited to develop more dangerous biological agents. Furthermore, unethical genetic alteration, such as conducting tests on individuals or animals without their consent, presents significant moral and legal dilemmas, including the potential for harm to subjects and the violation of their rights.

Artificial intelligence has exacerbated the risks associated with the misuse of biotechnology. Artificial intelligence can rapidly analyse genetic data, predict mutation occurrences, and automate laboratory procedures. These capabilities accelerate medical research; yet, they also facilitate the creation and utilisation of deleterious biological agents by malicious individuals, such as engineered pathogens or bioweapons, which pose significant threats to public health and safety. The collaboration of AI and biotechnology amplifies biological risks in both speed and scale, leading to potential threats such as the rapid development of harmful pathogens or bioweapons that could be exploited by malicious individuals.

In severe instances, the improper application of biotechnology may present an existential danger to humanity. Planned pandemics or uncontrollable biological agents could inundate healthcare systems, disrupt governments, and result in significant mortality. Biological threats pose a greater danger than conventional weapons due to their potential for silent and global dissemination, as they can spread rapidly through populations without immediate detection, leading to widespread illness and overwhelming healthcare resources.

Autonomous weapons use AI algorithms, sensors, and data processing to identify targets and make real-time decisions. Upon activation, autonomous weapons operate with minimal or no human supervision, in contrast to conventional weapons, which maintain human control over targeting and engagement. Examples include autonomous drones, robotic fighting systems, and collaborative swarm weaponry designed for coordinated attacks.

A significant ethical dilemma regarding autonomous weaponry is that they empower machines to make determinations concerning life and death. Excluding individuals from the decision-making process undermines moral accountability and human dignity. It is ethically indefensible to permit machines to determine life and death, as they lack empathy, moral discernment, and situational comprehension.

**The accountability gap**

The accountability deficit is an additional concern. Determining accountability for an autonomous weapon that perpetrates a war crime or inadvertently causes civilian casualties is challenging. Should it be the programmer, the military leader, the manufacturer, or the government? This ambiguity jeopardises international humanitarian law.

Contemporary international humanitarian rules are based on the assumption of human agency in armed conflict. Autonomous weapons complicate the application of criteria such as distinction, proportionality, and military necessity. The ability of AI systems to consistently distinguish between civilians and combatants remains highly questionable, hence elevating the potential of unlawful murders, which raises significant ethical concerns regarding the use of autonomous weapons in warfare.

Their velocity and autonomy could inadvertently exacerbate situations, as machines may react more swiftly than humans can intervene, potentially leading to unintended consequences such as escalated conflicts or accidents that humans cannot control in time. Furthermore, if a significant number of individuals begin utilising them, it may initiate a worldwide arms race, with both state and non-state entities striving to employ increasingly autonomous systems, potentially leading to heightened global tensions and conflicts over technological supremacy.

Autonomous weapons may pose a greater danger to humanity when integrated with cyber warfare and AI-driven decision-making. The widespread implementation of such systems may precipitate rapid and uncontrollable violence, so undermining global security and exacerbating conflicts beyond manageable levels, particularly as these systems could operate without human oversight and make decisions based on flawed algorithms or biased data.

**Safeguards**

Robust ethical principles in scientific inquiry

Regulation of dual-use biotechnology

Accountable incorporation of artificial intelligence in the life sciences

Global collaboration and openness

Enhancing international biosecurity legislation

**Counterarguments to Existential Risk**

Critics of the existential threat narrative argue that these concerns are overstated, speculative, or distract from more urgent matters.

- AGI is Far (or Impossible) It denotes the perspective that artificial general intelligence—AI possessing human-level, adaptable, and independent intelligence—is still distant from realisation, rendering existential or catastrophic anxieties premature. Certain scientists contend that modern AI models, including Large Language Models, are deficient in the requisite complex reasoning and autonomous learning capacities to attain Artificial General Intelligence (AGI). They contend that the trajectory toward authentic AGI is ambiguous and may necessitate decades, therefore affording humanity the chance to implement protections against potential risks associated with advanced AI systems, such as ethical concerns and unintended consequences.

**AGI faces several fundamental challenges:**

- Comprehension and Cognition: Currently, AI lacks a genuine understanding of the world. AI systems excel in pattern recognition; nevertheless, they lack the capacity for interdisciplinary thought that humans possess. Contemporary AI systems can identify patterns and generate outputs from data, but they do not possess true comprehension of their significance. A language model can compose essays; yet, it lacks comprehension of the real-world significance of its words. AI reasoning is constrained by its training data; it cannot extrapolate its knowledge to novel circumstances like humans can, which limits its ability to understand context and make informed decisions in unpredictable situations. Artificial Intelligence cannot function as Artificial General Intelligence without genuine comprehension and reasoning capabilities. The AI must possess the capability to strategise, draw conclusions, and independently resolve novel issues; mere task execution is insufficient. This constitutes a significant obstacle to attaining AGI.
- **Common Sense and Context:** Common sense constitutes the fundamental knowledge individuals possess regarding the functioning of many phenomena, facilitating their comprehension of circumstances. Context encompasses the particulars, objectives, and nuances that impart significance to real-life scenarios. Individuals frequently employ practical reasoning in their everyday activities, such as making decisions based on their experiences and the specific context of each situation. AI continues to struggle with context, ambiguity, and coherent thinking, which limits its ability to effectively understand and respond to the practical reasoning that individuals employ in their everyday activities. Individuals provide objectives, and the AI's responsibility is to determine the optimal methods for achieving those objectives while adhering to the constraints imposed by individuals, but without common sense, AI systems often fail to align their methods with the nuanced understanding of human values and

social contexts that are essential for effective decision-making. AI systems lack common sense, resulting in an absence of intuitive, everyday comprehension of the physical, social, and moral realms that individuals utilise for decision-making and judgement. Artificial intelligence cannot attain artificial general intelligence without possessing common sense and situational awareness. Highly sophisticated models can produce coherent outputs; nevertheless, they may also render illogical or hazardous conclusions when faced with novel or ambiguous scenarios, highlighting the necessity for common sense and situational awareness to navigate such complexities effectively.

- **Learning Efficiency:** Individuals can acquire the ability to perform complex tasks with merely a limited number of instances. Most AI systems require substantial amounts of data. Deep learning models and other AI systems require substantial data and extensive training to perform even basic tasks. A youngster can identify a new species from one or two images, whereas an AI model may require thousands of labelled instances. This inefficiency hinders AI's functionality across diverse real-world scenarios, particularly in situations where data is scarce or difficult to obtain. 'Data-dependent' signifies that AI systems require substantial pertinent data to learn, make decisions, and perform effectively. They are unable to function effectively or generalise without these extensive datasets. Individuals often fail to comprehend that contemporary AI systems lack a comprehensive, adaptable, human-like grasp of the universe. Rather than comprehending message, context, or intent, they operate by identifying patterns and correlations, which means they can only respond based on the data they have been trained on without true understanding of the underlying concepts. An AI system that is task-specific is designed and trained to do a singular task with high proficiency. It cannot apply its intellect to different activities or domains without undergoing retraining, which limits its versatility compared to more general AI systems that can adapt to various tasks. An AI system engineered for a singular purpose is capable of executing only one unique, predetermined task. It cannot acquire proficiency in additional tasks without significant effort or retraining. A primary reason AGI remains distant is its inadequate learning capabilities. Artificial intelligence cannot replicate human-like intellect or adapt to novel situations due to its inability to learn rapidly and transfer knowledge across diverse domains. Enhancing the efficiency of learning is crucial for bridging the divide between narrow AI and AGI, as it would enable AI systems to better generalise knowledge and apply it in various contexts, similar to human cognitive abilities.
- **Autonomy and Adaptability:** Autonomy refers to an AI system's capacity to make decisions and perform actions independently, without constant human assistance. 'Adaptability'

refers to the capacity to navigate novel circumstances, acquire knowledge from them, and effectively manage unforeseen challenges. AGI must possess the capability to independently plan, learn, and adapt to entirely novel situations, a proficiency that present AI systems do not consistently exhibit. Contemporary AI systems are predominantly designed for singular tasks and require human oversight, which limits their ability to function autonomously in dynamic environments and respond effectively to unforeseen challenges. They are unable to devise intricate strategies or adapt to significant changes, which limits their effectiveness in dynamic environments where flexibility and strategic planning are essential. An AI trained exclusively in one city may struggle to navigate another city with distinct traffic regulations or climatic conditions, which highlights the limitations of current AI systems in generalising knowledge and adapting to new environments. Genuine AGI, or Artificial General Intelligence, must possess the capability to independently learn and adapt across several domains, meaning it can understand and perform a wide range of tasks without human intervention. An AI system operating under defined objectives and constraints can solely execute tasks as instructed by humans. It is incapable of altering its objectives or going beyond its predefined constraints. Contemporary AI systems lack the capacity to establish, modify, or prioritise their objectives due to their inherent limitations. They are unable to modify their objectives based on the circumstances, which means they cannot adapt their goals in response to changing environments or new information. This indicates that AI systems currently do not establish their objectives.

- Ethical and Safety Concerns: Creating a system that is secure, congruent with human values, and resilient under all circumstances presents a considerable challenge.
- Manageability: Modern AI systems are primarily controllable and foreseeable. The notion of a completely autonomous, self-aware, and antagonistic AI is perhaps more consistent with science fiction than with actuality. Researchers continue to develop ways for AI alignment to ensure that sophisticated systems adhere to human ideals, such as ethical decision-making and transparency in their operations.

• Redirecting Focus from Present Issues: Excessive emphasis on abstract, future existential concerns diverts attention from the tangible, present impacts of AI, like algorithmic prejudice, employment displacement, misinformation, and the utilisation of AI for mass surveillance. These threats are tangible and require prompt action through the enactment of new legislation. Diverting attention is perilous, as the ramifications of existing AI applications are evident and imminent, whereas AGI remains merely a theoretical concept; therefore, it is essential to focus

on regulating current AI technologies to mitigate risks such as bias, privacy violations, and security threats. Addressing contemporary issues is both crucial and feasible, and it may also contribute to the long-term safety of AI research by ensuring that current technologies are developed responsibly and ethically, thereby preventing potential risks associated with their misuse.

The development of artificial intelligence raises important questions about its potential risks and benefits for humans. While AI has the capacity to greatly improve human life through progress in research, education, and technology, numerous philosophers warn that highly advanced AI systems may prove difficult to govern. If such systems are developed without sufficient safeguards, they may function in ways that oppose human values and interests.

**Knock your brain**

1 What is a major concern associated with superintelligent AI as an existential threat?

A) It may be slower than human intelligence.

B) It might pursue goals misaligned with human values.

C) It cannot process large datasets efficiently.

D) It requires human supervision for all tasks.

2 Why is the concept of “autonomy” in AI relevant to existential threats?

A) Autonomous AI is always safe.

B). Autonomous AI cannot learn new tasks

C) Autonomy reduces the need for computing resources.

D). Greater autonomy allows AI to act independently

3 What does AGI (Artificial General Intelligence) refer to?

A) AI that can perform any intellectual task a human can

B) AI specialized in one task, like chess or translation

C) AI that only stores data efficiently

D) AI that can only mimic human speech

4 Why is AGI considered “far” or a distant goal?

A) Because AI cannot play games efficiently

B) Because humans do not want AGI

C) Because current AI lacks understanding, reasoning, and common sense

D) Because computers are too expensive

5 How does the lack of autonomy in AI relate to AGI being far?

A) AI can already act completely independently

B) AI needs human supervision and cannot self-direct learning across domains

C) AI is too autonomous for safety

D) AI autonomy makes it easier to achieve AGI

**Unleash your potential**

1.Assess the necessity of imposing restrictions on AI development to mitigate potential hazards to humanity.

2.Is humanity capable of retaining control over highly powerful artificial intelligence systems? Engage in discourse.

3.Evaluate the benefits and possible risks of artificial intelligence for the future of human civilisation.